

<https://medium.com/@piyalidas.it/totp-time-based-one-time-password-what-if-the-phone-and-server-clocks-arent-exactly-6f999758f157>


Clock skew makes TOTP a distributed-systems problem disguised as authentication. A wider acceptance window improves reliability but increases replay opportunity. What combination of drift monitoring, one-time use, rate limits, and recovery keeps the user experience tolerant without making the window meaningless?




Yes — I'd frame the solution as **"small acceptance window + adaptive drift correction + atomic replay prevention,"** rather than simply making the TOTP window larger.

RFC 6238 itself recommends a 30-second default time step and says validators may accept prior time steps to account for transmission delay, but recommends **at most one time step** for that purpose because a larger window increases attack exposure.  RFC Editor +1

A practical production design





The key is to separate four concerns

Mechanism	Purpose	Recommended approach
Clock tolerance	Handle legitimate drift	<code>current ± 1 timestep</code>
Drift monitoring	Detect systematically wrong clocks	Record observed offset
One-time use	Stop replay	Store consumed timestep/OTP atomically
Rate limiting	Stop guessing	Per-account + IP/device limits
Recovery	Handle lost/broken authenticator	Backup codes / secure re-enrollment


OWASP similarly recommends short OTP TTLs, single use, strict attempt limits, and invalidating the OTP after successful verification. [🔗 OWASP Cheat S...](#)

Don't dynamically widen the window

This is the dangerous pattern:

```
graph TD; A[User fails] --> B[Try ±1]; B --> C[Fails]; C --> D[Try ±2]; D --> E[Fails]; E --> F[Try ±5]; F --> G["Let's accept it"];
```


Instead:



The RFC's resynchronization section explicitly discusses recording clock drift and compensating for it, rather than endlessly expanding the validation window.

Replay protection needs to be atomic

Suppose the phone is 20 seconds behind:



The important implementation detail is:

VERIFY + MARK_USED

should happen atomically.

For example, conceptually:

SQL



```
UPDATE totp_attempts
SET consumed = true
WHERE user_id = ?
      AND timestep = ?
      AND consumed = false;
```

Only the request that successfully changes the record gets to authenticate.

This matters because a **wider tolerance window isn't automatically catastrophic if every accepted timestep can only be consumed once**. But it still increases the period during which a stolen OTP can potentially be used, which is why RFC 6238 warns about larger windows.  RFC Editor

Drift monitoring can make the UX much better

Instead of storing only:

user → TOTP secret 

store something conceptually like:

```
user
├─ secret
├─ expected_period = 30s
├─ drift_steps = 0
├─ drift_confidence
└─ last_used_timestep 
```

If a user's OTP consistently validates at:

N-1
N-1
N-1 

you've learned:

client clock ≈ server clock - 30 seconds 

You can accommodate that **within a bounded policy**, while also flagging persistent drift.

But don't silently allow:

-1 → -2 → -3 → -4 → -5 → ... 

That turns "clock synchronization" into an attacker-controlled widening of the authentication window.

My recommended policy

For a typical web application:

TOTP period	30 seconds
Normal acceptance	current timestep
Tolerance	± 1 timestep
Maximum drift	bounded, e.g. 1-2 steps
Replay	atomic one-time consumption
Rate limit	per account + IP/device
Failed attempts	progressively throttled
Logging	never log OTP values
Monitoring	track drift distribution
Recovery	one-time backup codes
Severe drift	ask user to enable automatic time



The exact maximum drift should be a **risk-based product decision**, not a universal magic number.


And there's an important distinction:


Clock tolerance solves reliability. Replay protection solves stolen-code reuse. Rate limiting solves guessing. Recovery solves lockout.

Don't try to make one of these mechanisms compensate for another.

The distributed-systems perspective

Your Medium article's framing is actually quite good:





```
graph TD; A[ ] --> B[Bounded Acceptance]; B -- "+" --> C[Replay Prevention]; C -- "+" --> D[Rate Limiting]; D -- "+" --> E[Secure Recovery];
```

Bounded Acceptance
+
Replay Prevention
+
Rate Limiting
+
Secure Recovery

So I would summarize the design principle as:

Don't solve clock skew by making TOTP valid for longer. Solve it by making the validator tolerant to a small, measurable amount of drift while making every successful OTP consumption strictly one-time.

That gives you the sweet spot between **UX tolerance** and **authentication security**.

 RFC Editor +1



Is this conversation helpful so far?



ChatGPT can make mistakes. Check important info.